

AppNexus
Introduction to Python
Homework, Session 3

EXERCISES

Please note that sample output is sometimes split between pages in this document.
Also note that you must not use the CSV module for your exercises or homework! This is to give you a chance to use the "manual tools" before graduating to the "power tools".

Exercises related to Homework 3.1:

Ex. 3.1 open the FF-abbreviated.txt file, reading it line-by-line. print each line

Expected output:

```
19260701    0.09    0.22    0.30    0.009
19260702    0.44    0.35    0.08    0.009
19260706    0.17    0.26    0.37    0.009
...intermediate output omitted...
19280301    0.23    0.04    0.12    0.011
19280302    0.07    0.01    0.66    0.011
19280303    0.49    0.01    0.64    0.011
```

Ex. 3.2 building on the above program, set up a counter that is set to 0 before the loop begins, and counts 1 for each line in the file. print each line in the file, and at the end report the count.

Expected output bold added for emphasis, not part of output:

```
19260701    0.09    0.22    0.30    0.009
19260702    0.44    0.35    0.08    0.009
19260706    0.17    0.26    0.37    0.009
...intermediate output omitted...
19280301    0.23    0.04    0.12    0.011
19280302    0.07    0.01    0.66    0.011
19280303    0.49    0.01    0.64    0.011
```

Ex. 3.3 building on the above program, print the value of the counter in front of each line, so that each line is printed with its line number. (To print the number alongside of the line, you can use string formatting, concatenation with **str()**, or a comma, as in

```
print count, line
```

Expected output:

```
1 19260701      0.09      0.22      0.30      0.009
2 19260702      0.44      0.35      0.08      0.009
3 19260706      0.17      0.26      0.37      0.009
...intermediate output omitted...
24 19280301      0.23      0.04      0.12      0.011
25 19280302      0.07      0.01      0.66      0.011
26 19280303      0.49      0.01      0.64      0.011
26
```

Ex. 3.4 new program: open the FF-abbreviated file and print just the year from each line, so you see just a 4-digit year from each line. (hint: take a 4-digit slice of each line and instead of printing the line, print the slice)

Expected output:

```
1926
1926
1926
1926
1926
1926
1926
1926
1926
1926
1927
1927
1927
1927
1927
1927
1927
1927
1927
1928
1928
1928
1928
1928
1928
1928
1928
1928
1928
```

Ex. 3.5 building on the previous program, set a string variable to '1928'. comparing the string to the slice, print out only those lines where the strings are equivalent (i.e., the year for the line is '1928'). Keep in mind that the == operator works with numbers as well as strings

Expected output:

19280103	0.43	0.90	0.20	0.010
19280104	0.14	0.47	0.01	0.010
19280105	0.71	0.14	0.15	0.010
19280201	0.25	0.56	0.71	0.014
19280202	0.44	0.15	0.18	0.014
19280203	1.12	0.48	0.42	0.014
19280301	0.23	0.04	0.12	0.011
19280302	0.07	0.01	0.66	0.011
19280303	0.49	0.01	0.64	0.011

Ex. 3.6 building on the previous program, add the counter from one of the earlier programs and this time count just the lines that match the year '1928'. print the total count at the end. (Hint: make sure that the counter is incremented inside the if block, i.e., only if the year matches.)

Expected output:

19280103	0.43	0.90	0.20	0.010
19280104	0.14	0.47	0.01	0.010
19280105	0.71	0.14	0.15	0.010
19280201	0.25	0.56	0.71	0.014
19280202	0.44	0.15	0.18	0.014
19280203	1.12	0.48	0.42	0.014
19280301	0.23	0.04	0.12	0.011
19280302	0.07	0.01	0.66	0.011
19280303	0.49	0.01	0.64	0.011

Ex. 3.7 new program: open the FF-abbreviated file and split each line so the individual columns are separated into a list. print the list from each line. (Hint: **str.split()** without any argument splits on whitespace.)

Expected output:

```
['19260701', '0.09', '0.22', '0.30', '0.009']
['19260702', '0.44', '0.35', '0.08', '0.009']
['19260706', '0.17', '0.26', '0.37', '0.009']
['19260802', '0.82', '0.21', '0.01', '0.010']
['19260803', '0.46', '0.39', '0.38', '0.010']
['19260804', '0.35', '0.15', '0.32', '0.010']
['19260901', '0.54', '0.41', '0.08', '0.010']
['19260902', '0.04', '0.06', '0.23', '0.010']
['19260903', '0.48', '0.34', '0.09', '0.010']
['19270103', '0.97', '0.21', '0.24', '0.010']
['19270104', '0.30', '0.15', '0.73', '0.010']
['19270201', '0.00', '0.56', '1.09', '0.012']
['19270202', '0.72', '0.23', '0.18', '0.012']
['19270203', '0.17', '0.22', '0.08', '0.012']
['19270301', '0.38', '0.07', '0.57', '0.011']
['19270302', '1.12', '0.10', '0.22', '0.011']
['19270303', '1.01', '0.11', '0.04', '0.011']
['19280103', '0.43', '0.90', '0.20', '0.010']
['19280104', '0.14', '0.47', '0.01', '0.010']
['19280105', '0.71', '0.14', '0.15', '0.010']
['19280201', '0.25', '0.56', '0.71', '0.014']
['19280202', '0.44', '0.15', '0.18', '0.014']
['19280203', '1.12', '0.48', '0.42', '0.014']
['19280301', '0.23', '0.04', '0.12', '0.011']
['19280302', '0.07', '0.01', '0.66', '0.011']
['19280303', '0.49', '0.01', '0.64', '0.011']
```

Ex. 3.8 building on the previous program, instead of printing the entire list, print just the 1st column (the year-month-day) of each line. (Hint: since **str.split()** returns a list, you can easily print just one column from that list by using a subscript (i.e., index number inside square brackets after the list name).

Expected output:

```
19260701
19260702
19260706
19260802
19260803
19260804
19260901
19260902
19260903
19270103
19270104
19270201
19270202
19270203
19270301
```

```
19270302
19270303
19280103
19280104
19280105
19280201
19280202
19280203
19280301
19280302
19280303
```

Ex. 3.9 adjusting the previous program, print the the 2nd column instead of the 1st column.

Expected output:

```
0.09
0.44
0.17
0.82
0.46
0.35
0.54
0.04
0.48
0.97
0.30
0.00
0.72
0.17
0.38
1.12
1.01
0.43
0.14
0.71
0.25
0.44
1.12
0.23
0.07
0.49
```

Ex. 3.10 building on the previous program, now add the 'year selection' functionality from earlier and print only the 2nd column values whose lines match the year '1928'.
Note on efficiency: when adding in this functionality, you should make the line splitting and column selecting happen *only if* the year from the line is 1928. This means that the loop block will start with the slicing, then the 'if' test asking if the slice is equal to 1928, then inside that 'if', splitting the line, selecting the 2nd column, and printing the 2nd column value. The reason we want to follow this order is because we want the program to do as little work as possible: there's no point in splitting the line or selecting the value if the year doesn't match -- we'll be ignoring those lines anyway.

Expected output:

```
0.43
0.14
0.71
0.25
0.44
1.12
0.23
0.07
0.49
```

Ex. 3.11 building on the previous program, convert each of the 2nd column values to a float, and multiply that value * 2. Print the column value and then the doubled value on the same line (using a comma between them in a print statement is probably the easiest way to print a number and string together).

Expected output:

```
0.43 0.86
0.14 0.28
0.71 1.42
0.25 0.5
0.44 0.88
1.12 2.24
0.23 0.46
0.07 0.14
0.49 0.98
```

Ex. 3.12 building on the previous program, add in the counter, but increment the counter only if the year is 1928, so you're only counting the 1928 lines. print each count number, the float value, and the doubled float value on the same line.

Expected output:

```
1 0.43 0.86
2 0.14 0.28
3 0.71 1.42
4 0.25 0.5
5 0.44 0.88
6 1.12 2.24
7 0.23 0.46
8 0.07 0.14
9 0.49 0.98
```

Ex. 3.13 start a new program based on the logic of the previous one: create a sum of float values. before the loop begins, initialize a "floatsum" variable to 0. then looping through the data, if the year is 1928, select out the 2nd column (the 1st column of float values), convert it to a float, and add it to the floatsum variable. report the sum at the end.

Expected output:

3.88

Exercises related to Homework 3.2:

Ex. 3.14 Open the **pyku.txt** file and **file read()** the entire file into a string. Print the length of this string using **len()**

Expected output:

80

Ex. 3.15 Building on the previous program, count the number of times the word **spam** occurs in the file. (Hint: **file read()** the file into a string and use the **str count()** method.)

Expected output:

4

Ex. 3.16 Open the **pyku.txt** file and **file read()** it into a string. Use **str splitlines()** to split the string into a list of lines. Print the type of the list returned from **splitlines()**, then print the first and last line from the file using list subscripts.

Expected output:

```
<type 'list'>
We're out of gouda.
Spam, spam, spam, spam, spam.
```

Ex. 3.17 Open a file and **file.read()** it into a string. Use **str.split()** to split the entire string into a list of words. Print the type of the list returned from **split()**, then print the first and last word from the file using list subscripts.

Expected output:

```
<type 'list'>
We're
spam.
```

Ex. 3.18 Starting with this list:

```
var = ['a', 'b', 'c', 'd']
```

print the length of the list. (Hint: do not use a loop and a counter. Use **len()**)

Expected output:

4

HOMEWORK

Feel free to refer to the [discussion document](#) for further detail on how to proceed (including a code outline).

Also note that you must not use the CSV module for your exercises or homework! This is to give you a chance to use the "manual tools" before graduating to the "power tools".

- 3.1** Calculate the average Mkt-RF value (1st column of floats) for a given year. Use **raw_input()** to take user input -- a 4-digit year. Reject the input with a polite/rude message if it is not all digits. Looping through **Fama-French_data.txt** file, calculate the average **MktRF** value (the 1st column of floating-point values, i.e. the 2nd column in the file) for that year.

Note: it will be much easier to use **FF_abbreviated.txt** for testing, as you can mentally calculate the correct amount and check it against your program's output. Once this is confirmed, you can then run the program against the **Fama-French_data.txt** file and compare your output to the sample output below.

Sample program runs:

```
$ python 3-1.py
please enter a 4-digit year:  hello
sorry, that was bad input
please enter a 4-digit year:  what's wrong?
sorry, that was bad input
please enter a 4-digit year:  1990
count 253, sum -12.77, avg -0.0504743083004
```

```
$ python 3-1.py
please enter a 4-digit year:  1927
count 301, sum 26.26, avg 0.0872425249169
```

```
$ python 3-1.py
please enter a 4-digit year:  1945
count 286, sum 32.55, avg 0.113811188811
```


3.2 **wc** emulation: **wc** is a unix utility program that counts the number of lines, words and characters in a given file. In the below example command run on a file in the Unix os, **wc** tells us that the **FF_abbreviated.txt** file contains 25 lines, 130 words and 1065 characters (note that your count may be off by 1 or 2 lines, or up to 20 characters -- don't sweat such differences):

```
$ wc FF_abbreviated.txt
    25      130     1065 FF_abbreviated.txt
```

Our script will do the same work: reading from file **sawyer.txt** (found in the **source data** directory linked from the class website), print the number of lines, words and characters in the file. Note: when counting characters, include spaces and newlines as well.

Challenge 1: please do not open and read the file more than once.

Challenge 2: you will be tempted to loop and count to get your answer, **but you are challenged get these counts without looping**. See if you can get each count without actually using a loop, but by using **len()** on various "slice and dices" in the data. Do this without opening the file more than once (hint: **read()** will read the file into a string; **split()** will split a string into words; **splitlines()** will split a string on the newlines!).

Extra credit: attempt to mirror the format of **wc** by right-justifying each value within an 8-character width (use **str.rjust()**).

```
please enter a filename: sawyer.txt
    20      270     1435 sawyer.txt
```

Special note: we have seen anomalies that stem from varying line endings on the Windows and Mac platforms. You may get a count that is off mine by 1 or 2 lines, or up to 20 characters. If so, don't worry. It just needs to be within 20.

SUPPLEMENTARY (EXTRA CREDIT) EXERCISE

3.3 Write a Python program that can download CSV data from the internet and parse with the **CSV** module: real-world data is often more complex than the data we are working with, and supplementary exercises like these may be more complicated and/or time consuming than the "garden" assignments that are required. For these I will not work out solutions, so please enjoy the adventure and send me your questions.

CSV files need to be parsed with the **CSV** module, which takes into account some of these complexities (for example, including the delimiter in the data - what if you wanted your data to include a comma? Please see the new section titled "Modules: Internet Data" for details on using the **CSV** and **urllib2** modules.

Select one of the sources listed at the link below - pick something interesting, and something you feel you might be able to learn from based on the (albeit simple) techniques we've described.

https://catalog.data.gov/dataset?res_format=CSV&page=1

The yellow [CSV] button under each source is actually a link to CSV data (in most cases), but we don't want to download the data to a browser, so don't click the button. When you've found data you'd like to analyze, *right-click* this yellow [CSV] button and select "Copy Link Location". (If you don't see that choice, move your cursor on and off, and try again.) Copy this link to a plain text file for use in your program.

For example, here's the link copied from the very first data source on that page (the Consumer Complaint Database):

<https://data.consumerfinance.gov/api/views/s6ew-h6mp/rows.csv?accessType=DOWNLOAD>

Now use your selected link (not necessarily the one above) with the **urllib2** module to download the data of your choice from the website.

Remember that the handle you get back from **urllib.urlopen()** is readable in the same way a file is readable. Read the data and print the first few lines to see what you've got:

```
# my own urllib.urlopen() read handle I named rh:  
lines = rh.readlines()  
print lines[0:4]
```

Then you can decide which values you'd like to analyze (at this point, we can only do sums and averages).

Send me your questions. Enjoy!